

## 前言

redis 是我们目前大规模使用的缓存中间件，由于它强大高效而又便捷的功能，得到了广泛的使用。现在的2.x的稳定版本是2.8.19,也是我们项目中普遍用到的版本。redis在年初发布了3.0.0,官方支持了redis cluster,也就是集群。至此结束了redis 没有官方集群的时代，之前我们用redis cluster用的最多的应该是twitter 发布的Twemproxy(<https://github.com/twitter/twemproxy>)

还有就是豌豆荚开发的codis(<https://github.com/wandoulabs/codis>)

这2个我会在后续去使用它们。下面的文字，我尽量用通俗易懂的方式来阐述。

### redis cluster 理论知识

截止写这篇文章前，redis 3.x的最新版本是3.0.5。今天的学习和实践就是用这个版本。

#### Redis Cluster设计要点

redis cluster在设计的时候，就考虑到了去中心化，去中间件，也就是说，集群中的每个节点都是平等的关系，都是对等的，每个节点都保存各自的数据和整个集群的状态。每个节点都和其他所有节点连接，而且这些连接保持活跃，这样就保证了我们只需要连接集群中的任意一个节点，就可以获取到其他节点的数据。

那么redis 是如何合理分配这些节点和数据的呢？

Redis 集群没有并使用传统的一致性哈希来分配数据，而是采用另外一种叫做哈希槽（hash slot）的方式来分配的。redis cluster 默认分配了 16384 个slot，当我们set一个key 时，会用CRC16算法来取模得到所属的slot，然后将这个key 分到哈希槽区间的节点上，具体算法就是： $CRC16(key) \% 16384$ 。

注意的是：必须要3个以后的主节点，否则在创建集群时会失败，我们在后续会实践到。

所以，我们假设现在有3个节点已经组成了集群，分别是：A, B, C 三个节点，它们可以是一台机器上的三个端口，也可以是三台不同的服务器。那么，采用哈希槽（hash slot）的方式来分配16384个slot 的话，它们三个节点分别承担的slot 区间是：

节点A覆盖0 - 5460;节点B覆盖5461 - 10922;节点C覆盖10923 - 16383。

如下图所示：



那么，现在我想设置一个key ,比如叫my\_name:

```
set my_name yangyi
```

按照redis cluster的哈希槽算法： $CRC16(\&\#39;my\_name\&\#39;)\%16384 = 2412$ 。那么就会把这个key 的存储分配到 A 上了。


同样，当我连接(A,B,C)任何一个节点想获取my\_name这个key时，也会这样的算法，然后内部跳转到B节点上获取数据。

这种哈希槽的分配方式有好也有坏，好处就是很清晰，比如我想新增一个节点D，redis cluster的这种做法是从各个节点的前面各拿取一部分slot到D上，我会在接下来的实践中实验。大致就会变成这样：

节点A覆盖1365-5460节点B覆盖6827-10922节点C覆盖12288-16383节点D覆盖0-1364,5461-6826,10923-12287

同样删除一个节点也是类似，移动完成后就可以删除这个节点了。

所以redis cluster 就是这样的一个形状：



#### Redis Cluster主从模式

redis cluster 为了保证数据的高可用性，加入了主从模式，一个主节点对应一个或多个从节点，主节点提供数据存取，从节点则是从主节点拉取数据备份，当这个主节点挂掉后，就会有这个从节点选取一个来充当主节点，从而保证集群不会挂掉。


上面那个例子里，集群有ABC三个主节点，如果这3个节点都没有加入从节点，如果B挂掉了，我们就无法访问整个集群了。A和C的slot也无法访问。

所以我们在集群建立的时候，一定要为每个主节点都添加了从节点。比如像这样，集群包含主节点A、B、C，以及从节点A1、B1、C1，那么即使B挂掉系统也可以继续正常工作。

B1节点替代了B节点，所以Redis集群将会选择B1节点作为新的主节点，集群将会继续正确地提供服务。当B重新开启后，它就会变成B1的从节点。

不过需要注意，如果节点B和B1同时挂了，Redis集群就无法继续正确地提供服务了。

流程图所示：



#### redis cluster 动手实践

按照官方的教程，全程再演示一次，官方教程是在一台机器上启动6个节点，3个当主，3个当从(<http://redis.io/topics/cluster-tutorial>):  
先下载官方的redis 版本(3.0.5) :<http://download.redis.io/releases/redis-3.0.5.tar.gz>

下载不了，请自行翻墙。我这次是在centos6.5上演示，用的是root 账户。

如果之前已经下载了redis的 2.x版本，只需要将/usr/local/bin/redis-\*这几个命令先删除即可。

## 1.解压

```
1. [root@web3 ~]# tar zxvf redis-3.0.5.tar.gz
```

## 2.安装

```
1. [root@web3 ~]# cd redis-3.0.5
```

```
2. [root@web3 ~]# make && make install
```

## 3.将redis-trib.rb 复制到/usr/local/bin

```
1. [root@web3 redis-3.0.5]# cd src/
```

```
2. [root@web3 src]# cp redis-trib.rb /usr/local/bin
```

## 4.开始集群搭建

```
1. [root@web3 redis-3.0.5]# vi redis.conf
```

2. #修改以下地方

3. port 7000

4. cluster-enabled yes

5. cluster-config-file nodes.conf

6. cluster-node-timeout 5000

7. appendonly yes

新建6个节点：

```
1. [root@web3 redis-3.0.5]# mkdir -p /usr/local/cluster-test
```

```
2. [root@web3 redis-3.0.5]# cd /usr/local/cluster-test/
```

```
3. [root@web3 cluster-test]# mkdir 7000
```

```
4. [root@web3 cluster-test]# mkdir 7001
```

```
5. [root@web3 cluster-test]# mkdir 7002
```

```
6. [root@web3 cluster-test]# mkdir 7003
```

```
7. [root@web3 cluster-test]# mkdir 7004
```

```
8. [root@web3 cluster-test]# mkdir 7005
```

将redis.conf 分别拷贝到这6个文件夹中,并修改成对应的端口号

1. #拷贝配置文件

```
2. [root@web3 cluster-test]# cp /root/redis-3.0.5/redis.conf /usr/local/cluster-test/7000
```

```
3. [root@web3 cluster-test]# cp /root/redis-3.0.5/redis.conf /usr/local/cluster-test/7001
```

```
4. [root@web3 cluster-test]# cp /root/redis-3.0.5/redis.conf /usr/local/cluster-test/7002
```

```
5. [root@web3 cluster-test]# cp /root/redis-3.0.5/redis.conf /usr/local/cluster-test/7003
```

```
6. [root@web3 cluster-test]# cp /root/redis-3.0.5/redis.conf /usr/local/cluster-test/7004
```

```
7. [root@web3 cluster-test]# cp /root/redis-3.0.5/redis.conf /usr/local/cluster-test/7005
```

8. # 修改端口号

```
9. root@web3 cluster-test]# sed -i "s/7000/7001/g" /usr/local/cluster-test/7001/redis.conf
```

```
0. [root@web3 cluster-test]# sed -i "s/7000/7002/g" /usr/local/cluster-test/7002/redis.conf
```

```
1. [root@web3 cluster-test]# sed -i "s/7000/7003/g" /usr/local/cluster-test/7003/redis.conf
```

```
2. [root@web3 cluster-test]# sed -i "s/7000/7004/g" /usr/local/cluster-test/7004/redis.conf
```

```
3. [root@web3 cluster-test]# sed -i "s/7000/7005/g" /usr/local/cluster-test/7005/redis.conf
```

分别，启动这6个节点

```
1. [root@web3 cluster-test]# cd /usr/local/cluster-test/7000/
```

```
2. [root@web3 7000]# redis-server redis.conf
```

```
3. [root@web3 7000]# cd ../7001
```

```
4. [root@web3 7001]# redis-server redis.conf
```

```
5. [root@web3 7001]# cd ../7002
```

```
6. [root@web3 7002]# redis-server redis.conf
```

```
7. [root@web3 7002]# cd ../7003
```

```
8. [root@web3 7003]# redis-server redis.conf
```

```
9. [root@web3 7003]# cd ../7004
```

```
0. [root@web3 7004]# redis-server redis.conf
```

```
1. [root@web3 7004]# cd ../7005
```

```
2. [root@web3 7005]# redis-server redis.conf
```

```
3. [root@web3 7005]#
```

查看6个节点的启动进程情况：

```
1. [root@web3 7005]# ps -ef|grep redis
```

```
2. root 11380 1 0 07:37 ? 00:00:00 redis-server *:7000 [cluster]
```

```
3. root 11384 1 0 07:37 ? 00:00:00 redis-server *:7001 [cluster]
```

```
4. root 11388 1 0 07:37 ? 00:00:00 redis-server *:7002 [cluster]
```

```
5. root 11392 1 0 07:37 ? 00:00:00 redis-server *:7003 [cluster]
```

```
6. root 11396 1 0 07:37 ? 00:00:00 redis-server *:7004 [cluster]
```

```
7. root 11400 1 0 07:37 ? 00:00:00 redis-server *:7005 [cluster]
```

```
8. root 11404 8259 0 07:38 pts/0 00:00:00 grep redis
```

将6个节点连在一起构招成集群

需要用到命令就是redis-trib.rb,这是官方的一个用ruby写的一个操作redis cluster的命令，所以，你的机器上需要安装ruby。我们先试一下这个命令：

```
1.redis-trib.rb create --replicas 1 127.0.0.1:7000 127.0.0.1:7001 127.0.0.1:7002 127.0.0.1:7003 127.0.0.1:7004 127.0.0.1:7005
```

因为我们要新建集群, 所以这里使用create命令, --replicas 1参数表示为每个主节点创建一个从节点. 其他参数是实例的地址集合。

由于我机子上没安装ruby, 所以, 会报错:

```
1./usr/bin/env: ruby: No such file or directory
```

那先安装ruby和rubygems:

```
1.[root@web3 7005]# yum install ruby ruby-devel rubygems rpm-build
```

然后, 再执行一次, 发现还是报错:

```
1.[root@web3 7005]# redis-trib.rb create --replicas 1 127.0.0.1:7000 127.0.0.1:7001 127.0.0.1:7002 127.0.0.1:7003 127.0.0.1:7004 127.0.0.1:7005
```

```
2./usr/lib/ruby/site_ruby/1.8/rubygems/custom_require.rb:31:in `gem_original_require': no such file to load -- redis (LoadError)
```

```
3.from /usr/lib/ruby/site_ruby/1.8/rubygems/custom_require.rb:31:in `require';
```

```
4.from /usr/local/bin/redis-trib.rb:25
```

```
5.[root@web3 7005]#
```

原来是ruby和redis的连接没装好:

```
1.[root@web3 7005]# gem install redis
```

```
2.
```

```
3.Successfully installed redis-3.2.1
```

```
4.1 gem installed
```

```
5.Installing ri documentation for redis-3.2.1...
```

```
6.Installing RDoc documentation for redis-3.2.1...
```

好了。最重要的时刻来临了, 胜败成举在此了, 啊啊啊啊啊啊, 好激动:

```
1.[root@web3 7005]# redis-trib.rb create --replicas 1 127.0.0.1:7000 127.0.0.1:7001 127.0.0.1:7002 127.0.0.1:7003 127.0.0.1:7004 127.0.0.1:7005
```

```
2.>>> Creating cluster
```

```
3.Connecting to node 127.0.0.1:7000: OK
```

```
4.Connecting to node 127.0.0.1:7001: OK
```

```
5.Connecting to node 127.0.0.1:7002: OK
```

```
6.Connecting to node 127.0.0.1:7003: OK
```

```
7.Connecting to node 127.0.0.1:7004: OK
```

```
8.Connecting to node 127.0.0.1:7005: OK
```

```
9.>>> Performing hash slots allocation on 6 nodes...
```

```
0.Using 3 masters:
```

```
1.127.0.0.1:7000
```

```
2.127.0.0.1:7001
```

```
3.127.0.0.1:7002
```

```
4.Adding replica 127.0.0.1:7003 to 127.0.0.1:7000
```

```
5.Adding replica 127.0.0.1:7004 to 127.0.0.1:7001
```

```
6.Adding replica 127.0.0.1:7005 to 127.0.0.1:7002
```

```
7.M: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
```

```
8.slots:0-5460 (5461 slots) master
```

```
9.M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
```

```
10.slots:5461-10922 (5462 slots) master
```

```
11.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
```

```
12.slots:10923-16383 (5461 slots) master
```

```
13.S: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
```

```
14.replicates 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4
```

```
15.S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
```

```
16.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
```

```
17.S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
```

```
18.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
```

```
19.Can I set the above configuration? (type &#39;yes&#39; to accept): yes
```

```
20.>>> Nodes configuration updated
```

```
21.>>> Assign a different config epoch to each node
```

```
22.>>> Sending CLUSTER MEET messages to join the cluster
```

```
23.Waiting for the cluster to join...
```

```
24.>>> Performing Cluster Check (using node 127.0.0.1:7000)
```

```
25.M: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
```

```
26.slots:0-5460 (5461 slots) master
```

```
27.M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
```

```
28.slots:5461-10922 (5462 slots) master
```

```
29.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
```

```
30.slots:10923-16383 (5461 slots) master
```

```
31.M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
```

```
32.slots: (0 slots) master
```

```
33.replicates 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4
```

```
34.M: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
```

```
35.slots: (0 slots) master
```

```
36.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
```

```
37.M: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
```

```
38.slots: (0 slots) master
```

```
39.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
```

```
40.[OK] All nodes agree about slots configuration.
```

```
41.>>> Check for open slots...
```

```
i2.>>> Check slots coverage...
i3. [OK] All 16384 slots covered.
    哈哈哈哈哈。终于他妈的成功了!!!!
    Redis-trib会提示你做了什么配置, 输入yes接受. 集群就被配置和加入了, 意思是, 实例会经过互相交流后启动。
    测试集群的状态:
1. [root@web3 7000]# redis-trib.rb check 127.0.0.1:7000
2. Connecting to node 127.0.0.1:7000: OK
3. Connecting to node 127.0.0.1:7002: OK
4. Connecting to node 127.0.0.1:7003: OK
5. Connecting to node 127.0.0.1:7005: OK
6. Connecting to node 127.0.0.1:7001: OK
7. Connecting to node 127.0.0.1:7004: OK
8.>>> Performing Cluster Check (using node 127.0.0.1:7000)
9. M: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
0. slots:0-5460 (5461 slots) master
1. 1 additional replica(s)
2. M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
3. slots:10923-16383 (5461 slots) master
4. 1 additional replica(s)
5. S: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
6. slots: (0 slots) slave
7. replicates 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4
8. S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
9. slots: (0 slots) slave
0. replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
!1. M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
!2. slots:5461-10922 (5462 slots) master
!3. 1 additional replica(s)
!4. S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
!5. slots: (0 slots) slave
!6. replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
!7. [OK] All nodes agree about slots configuration.
!8.>>> Check for open slots...
!9.>>> Check slots coverage...
!0. [OK] All 16384 slots covered.
```

可以看到有3个主节点, 3个从节点。每个节点都是成功的连接状态。

3个主节点 [M] 是:

```
7000 (3707debcbe7be66d4a1968eaf3a5ffaf4308efa4)
7001 (cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c)
7002 (dfa0754c7854a874a6ebd2613b86140ad97701fc)
```

3个从节点 [S] 是:

```
7003 (d2237fdcfbba672de766b913d1186cebc6e1761)->7000
7004 (4b4aef8b48c427a3c903518339d53b6447c58b93)->7001
7005 (30858dbf483b61b9838d5c1f853a60beaa4e7afd) ->7002
```

## 5. 测试连接集群

刚才集群搭建成功了。按照redis cluster的特点, 它是去中心化, 每个节点都是对等的, 所以, 你连接哪个节点都可以获取和设置数据, 我们来试一下。

redis-cli是redis默认的客户工具, 启动时加上`-c`参数, 就可以连接到集群。

连接任意一个节点端口:

```
1. [root@web3 7000]# redis-cli -c -p 7000
2. 127.0.0.1:7000>
```

设置一个值:

```
1. 127.0.0.1:7000> set my_name yangyi
2. -> Redirected to slot [12803] located at 127.0.0.1:7002
3. OK
4. 127.0.0.1:7002> get my_name
5. "yangyi"
6. 127.0.0.1:7002>
```

前面理论知识我们知道了, 分配key的时候, 它会使用CRC16(&#39;my\_name&#39;)%16384算法, 来计算, 将这个key 放到哪个节点, 这里分配到了12803slot 就分配到了7002(10923-16383)这个节点上。

Redirected to slot [12803] located at 127.0.0.1:7002

redis cluster 采用的方式很直接, 它直接跳转到7002 节点了, 而不是还在自身的7000节点。

好, 现在我们连接7005这个从节点:

```
1. [root@web3 7000]# redis-cli -c -p 7005
2. 127.0.0.1:7005> get my_name
3. -> Redirected to slot [12803] located at 127.0.0.1:7002
4. "yangyi"
5. 127.0.0.1:7002>
```

我们同样是获取my\_name的值, 它同样也是跳转到了7002上。

我们再尝试一些其他的可以:

```
1. 127.0.0.1:7002> set age 123
```

```

2.-> Redirected to slot [741] located at 127.0.0.1:7000
3.OK
4.127.0.0.1:7000> set height 565
5.-> Redirected to slot [8223] located at 127.0.0.1:7001
6.OK
7.127.0.0.1:7001> set sex 1
8.-> Redirected to slot [2584] located at 127.0.0.1:7000
9.OK
0.127.0.0.1:7000> set home china
1.-> Redirected to slot [10814] located at 127.0.0.1:7001
2.OK
3.127.0.0.1:7001> set city shanghai
4.-> Redirected to slot [11479] located at 127.0.0.1:7002
5.OK
6.127.0.0.1:7002> set citylocate shanghaipudong
7.OK
8.127.0.0.1:7001> set wsdwxzx hhh
9.-> Redirected to slot [15487] located at 127.0.0.1:7002
0.OK
1.127.0.0.1:7002> zadd erew 333 rrr
2.-> Redirected to slot [10576] located at 127.0.0.1:7001
3.(integer) 1
4.127.0.0.1:7000> zrange erew 0 -1
5.-> Redirected to slot [10576] located at 127.0.0.1:7001
6.1) "rrr"
7.127.0.0.1:7001>

```

可以看出，数据都会在7000-7002 这3个主节点来跳转存粗。

## 6. 测试集群中的节点挂掉

上面我们建立来了一个集群。3个主节点 [7000-7002] 提供数据存粗和读取，3个从节点 [7003-7005] 则是负责把 [7000-7002] 的数据同步到自己的节点上来，我们来看一下 [7003-7005] 的appendonly.aof的内容。看看是不是不这样：

```

1.[root@web3 7005]# cd /usr/local/cluster-test/7003
2.[root@web3 7003]# vi appendonly.aof
3.
4.*2
5.$6
6.SELECT
7.$1
8.0
9.*3
0.$3
1.set
2.$3
3.age
4.$3
5.123
6.*3
7.$3
8.set
9.$3
0.sex
1.$1
2.1
3.*3
4.$3
5.set
6.$3
7.job
8.$3
9.php

```

我们看下，的确是从7000节点上同步过来的数据，7004，7005也是。

下面，我们先来模拟其中一台Master主服务器挂掉的情况，那就7000挂掉吧：

```

1.[root@web3 7003]# ps -ef|grep redis
2.root 11380 1 0 07:37 ? 00:00:03 redis-server *:7000 [cluster]
3.root 11384 1 0 07:37 ? 00:00:03 redis-server *:7001 [cluster]
4.root 11388 1 0 07:37 ? 00:00:03 redis-server *:7002 [cluster]
5.root 11392 1 0 07:37 ? 00:00:03 redis-server *:7003 [cluster]
6.root 11396 1 0 07:37 ? 00:00:04 redis-server *:7004 [cluster]
7.root 11400 1 0 07:37 ? 00:00:03 redis-server *:7005 [cluster]
8.[root@web3 7003]# kill 11380
9.[root@web3 7003]#

```

好，安装前面的理论，7000主节点挂掉了，那么这个时候，7000的从节点只有7003一个，肯定7003就会被选举称Master节点了：

```
1. [root@web3 7003]# redis-trib.rb check 127.0.0.1:7000
2. Connecting to node 127.0.0.1:7000: [ERR] Sorry, can't connect to node 127.0.0.1:7000
3. [root@web3 7003]# redis-trib.rb check 127.0.0.1:7001
4. Connecting to node 127.0.0.1:7001: OK
5. Connecting to node 127.0.0.1:7004: OK
6. Connecting to node 127.0.0.1:7005: OK
7. Connecting to node 127.0.0.1:7003: OK
8. Connecting to node 127.0.0.1:7002: OK
9. >>> Performing Cluster Check (using node 127.0.0.1:7001)
0. M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
   1.slots:5461-10922 (5462 slots) master
   2.1 additional replica(s)
3. S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
   4.slots: (0 slots) slave
   5.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
6. S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
   7.slots: (0 slots) slave
   8.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
9. M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
10.slots:0-5460 (5461 slots) master
11.0 additional replica(s)
12.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
13.slots:10923-16383 (5461 slots) master
14.1 additional replica(s)
15. [OK] All nodes agree about slots configuration.
16. >>> Check for open slots...
17. >>> Check slots coverage...
18. [OK] All 16384 slots covered.
19. [root@web3 7003]#
```

看了下，上面有了三个M节点了，果真，7003被选取成了替代7000成为主节点了。那我们来获取原先存在7000节点的数据：

```
1. [root@web3 7003]# redis-cli -c -p 7001
2. 127.0.0.1:7001> get sex
3. -> Redirected to slot [2584] located at 127.0.0.1:7003
4. "1"
5. 127.0.0.1:7003>
```

数据果真没有丢失，而是从7003上面获取了。

OK。我们再来模拟 7000节点重新启动了的情况，那么它还会自动加入到集群中吗？那么，7000这个节点上充当什么角色呢？我们试一下：  
重新启动 7000 节点：

```
1. [root@web3 7003]# cd ../7000
2. [root@web3 7000]# ll
3. total 56
4. -rw-r--r-- 1 root root 114 Oct 17 08:16 appendonly.aof
5. -rw-r--r-- 1 root root 43 Oct 17 08:37 dump.rdb
6. -rw-r--r-- 1 root root 745 Oct 17 08:00 nodes.conf
7. -rw-r--r-- 1 root root 41550 Oct 17 07:37 redis.conf
8. [root@web3 7000]# redis-server redis.conf
```

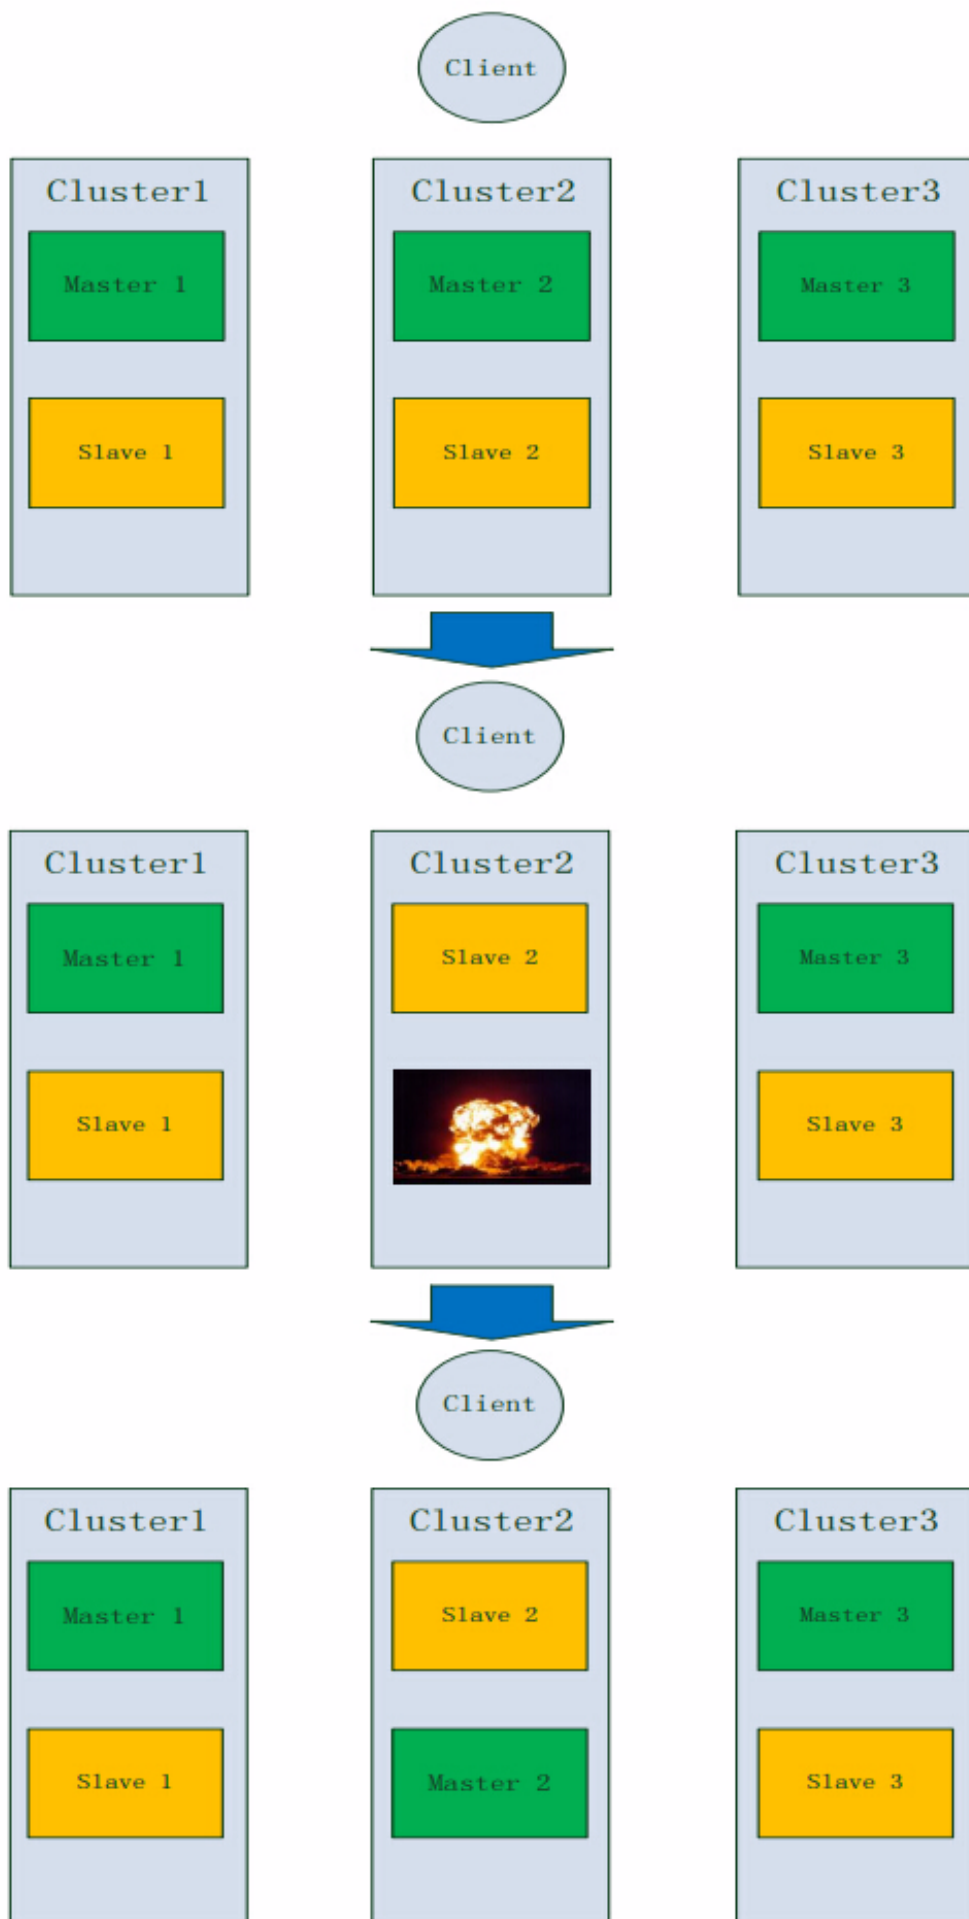
启动好了，现在，再来检查一下集群：

```
1. redis-trib.rb check 127.0.0.1:7000
1. [root@web3 7000]# redis-trib.rb check 127.0.0.1:7001
2. Connecting to node 127.0.0.1:7001: OK
3. Connecting to node 127.0.0.1:7004: OK
4. Connecting to node 127.0.0.1:7005: OK
5. Connecting to node 127.0.0.1:7003: OK
6. Connecting to node 127.0.0.1:7000: OK
7. Connecting to node 127.0.0.1:7002: OK
8. >>> Performing Cluster Check (using node 127.0.0.1:7001)
9. M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
   0.slots:5461-10922 (5462 slots) master
   1.1 additional replica(s)
2. S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
   3.slots: (0 slots) slave
   4.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
5. S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
   6.slots: (0 slots) slave
   7.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
8. M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
   9.slots:0-5460 (5461 slots) master
10.1 additional replica(s)
11.S: 3707debcb7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
12.slots: (0 slots) slave
```



```
'3.replicates d2237fdcfbba672de766b913d1186cebc6e1761
'4.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
'5.slots:10923-16383 (5461 slots) master
'6.1 additional replica(s)
'7.[OK] All nodes agree about slots configuration.
'8.>>> Check for open slots...
'9.>>> Check slots coverage...
'0.[OK] All 16384 slots covered.
```

你看，7000节点启动起来了，它却作为了70003的从节点了。验证了之前的这张图：



一定要保证有3个master 节点，不然，集群就挂掉了。

## 7. 集群中新加入节点

我们再来测试一下，新加入一个节点，分2种情况，1是作为主节点，2是作为一个节点的从节点。我们分别来试一下：

## 1. 新建一个7006节点 作为一个新的主节点加入：

首先就是新建一个 7006的文件夹和redis.conf:

```
1. [root@web3 cluster-test]# cd /usr/local/cluster-test/
2. [root@web3 cluster-test]# mkdir 7006
3. [root@web3 cluster-test]# cp 7005/redis.conf 7006/redis.conf
4. #修改端口
5. [root@web3 cluster-test]# sed -i "s/7005/7006/g" /usr/local/cluster-test/7006/redis.conf
启动 7006
```

```
1. [root@web3 7006]# redis-server redis.conf
2. [root@web3 7006]# ps -ef|grep redis
3. root 11384 1 0 07:37 ? 00:00:05 redis-server *:7001 [cluster]
4. root 11388 1 0 07:37 ? 00:00:05 redis-server *:7002 [cluster]
5. root 11392 1 0 07:37 ? 00:00:05 redis-server *:7003 [cluster]
6. root 11396 1 0 07:37 ? 00:00:06 redis-server *:7004 [cluster]
7. root 11400 1 0 07:37 ? 00:00:05 redis-server *:7005 [cluster]
8. root 12100 1 0 08:42 ? 00:00:01 redis-server *:7000 [cluster]
9. root 12132 1 0 09:09 ? 00:00:00 redis-server *:7006 [cluster]
0. root 12136 8259 0 09:10 pts/0 00:00:00 grep redis
```

ok,7006 端口已经启动好了，并且进程也存在了，下面就是加入到集群中，也是得用到redis-trib.rb命令：

```
1. redis-trib.rb add-node 127.0.0.1:7006 127.0.0.1:7000
add-node是加入指令，127.0.0.1:7006表示新加入的节点，127.0.0.1:7000表示加入的集群的一个节点，用来辨识是哪个集群，理论上哪个都可以。
```

```
1. [root@web3 7006]# redis-trib.rb add-node 127.0.0.1:7006 127.0.0.1:7000
```

```
2. >>> Adding node 127.0.0.1:7006 to cluster 127.0.0.1:7000
```

```
3. Connecting to node 127.0.0.1:7000: OK
```

```
4. Connecting to node 127.0.0.1:7004: OK
```

```
5. Connecting to node 127.0.0.1:7005: OK
```

```
6. Connecting to node 127.0.0.1:7003: OK
```

```
7. Connecting to node 127.0.0.1:7001: OK
```

```
8. Connecting to node 127.0.0.1:7002: OK
```

```
9. >>> Performing Cluster Check (using node 127.0.0.1:7000)
```

```
0. S: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
```

```
1. slots: (0 slots) slave
```

```
2. replicates d2237fdcfbba672de766b913d1186cebc6e1761
```

```
3. S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
```

```
4. slots: (0 slots) slave
```

```
5. replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
```

```
6. S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
```

```
7. slots: (0 slots) slave
```

```
8. replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
```

```
9. M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
```

```
10. slots:0-5460 (5461 slots) master
```

```
11. 1 additional replica(s)
```

```
12. M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
```

```
13. slots:5461-10922 (5462 slots) master
```

```
14. 1 additional replica(s)
```

```
15. M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
```

```
16. slots:10923-16383 (5461 slots) master
```

```
17. 1 additional replica(s)
```

```
18. [OK] All nodes agree about slots configuration.
```

```
19. >>> Check for open slots...
```

```
20. >>> Check slots coverage...
```

```
21. [OK] All 16384 slots covered.
```

```
22. Connecting to node 127.0.0.1:7006: OK
```

```
23. >>> Send CLUSTER MEET to node 127.0.0.1:7006 to make it join the cluster.
```

```
24. [OK] New node added correctly.
```

这个命令执行完成之后，它顺便检查来其他的6个节点都是成功的，最后几句话：

```
1. Connecting to node 127.0.0.1:7006: OK
```

```
2. >>> Send CLUSTER MEET to node 127.0.0.1:7006 to make it join the cluster.
```

```
3. [OK] New node added correctly.
```

表示新的节点连接成功了，而且也已经加入到集群了，我们再来检查一下：

```
1. [root@web3 7006]# redis-trib.rb check 127.0.0.1:7006
```

```
2. Connecting to node 127.0.0.1:7006: OK
```

```
3. Connecting to node 127.0.0.1:7001: OK
```

```
4. Connecting to node 127.0.0.1:7004: OK
```

```
5. Connecting to node 127.0.0.1:7000: OK
```

```
6. Connecting to node 127.0.0.1:7002: OK
```

```
7. Connecting to node 127.0.0.1:7005: OK
```

```
8. Connecting to node 127.0.0.1:7003: OK
```

```
9. >>> Performing Cluster Check (using node 127.0.0.1:7006)
```

```
0. M: efc3131fbdc6cf929720e0e0f7136cae85657481 127.0.0.1:7006
```

```
1.slots: (0 slots) master
2.0 additional replica(s)
3.M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
4.slots:5461-10922 (5462 slots) master
5.1 additional replica(s)
6.S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
7.slots: (0 slots) slave
8.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
9.S: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
10.slots: (0 slots) slave
11.replicates d2237fdcfbba672de766b913d1186cebc6e1761
12.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
13.slots:10923-16383 (5461 slots) master
14.1 additional replica(s)
15.S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
16.slots: (0 slots) slave
17.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
18.M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
19.slots:0-5460 (5461 slots) master
20.1 additional replica(s)
21.[OK] All nodes agree about slots configuration.
22.>>> Check for open slots...
23.>>> Check slots coverage...
24.[OK] All 16384 slots covered.
```

可以看到，7006 已经成为了主节点。

我们也可以连接到客户端后，来看这个集群节点情况：

```
1.[root@web3 7006]# redis-cli -c -p 7006
2.127.0.0.1:7006> cluster nodes
3.cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001 master - 0 1445073797986 2 connected 5461-10922
4.4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004 slave cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 0 1445073799497 2 connected
5.efc3131fbdc6cf929720e0ef7136cae85657481 127.0.0.1:7006 myself,master - 0 0 0 connected
6.3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000 slave d2237fdcfbba672de766b913d1186cebc6e1761 0 1445073797482 7 connected
7.dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002 master - 0 1445073798489 3 connected 10923-16383
8.30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005 slave dfa0754c7854a874a6ebd2613b86140ad97701fc 0 1445073798993 3 connected
9.d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003 master - 0 1445073799498 7 connected 0-5460
0.127.0.0.1:7006>
```

可以看到 有7个节点。7006 也作为master节点，但是，但是，你看看后面的：

```
1.efc3131fbdc6cf929720e0ef7136cae85657481 127.0.0.1:7006
2.slots: (0 slots) master
```

卧槽，什么情况。o slots也就是说，虽然它现在是主节点，但是，却没有分配任何节点给它，也就是它现在还不负责数据存取。那加上有毛用啊！！！！

看来，redis cluster 不是在新加节点的时候帮我们做好了迁移工作，需要我们手动对集群进行重新分片迁移，也是这个命令：

```
redis-trib.rb reshard 127.0.0.1:7000
```

这个命令是用来迁移slot节点的，后面的127.0.0.1:7000是表示是哪个集群，端口填 [7000-7006] 都可以，我们运行下：

```
1.[root@web3 7006]# redis-trib.rb reshard 127.0.0.1:7000
2.Connecting to node 127.0.0.1:7006: OK
3.Connecting to node 127.0.0.1:7001: OK
4.Connecting to node 127.0.0.1:7004: OK
5.Connecting to node 127.0.0.1:7000: OK
6.Connecting to node 127.0.0.1:7002: OK
7.Connecting to node 127.0.0.1:7005: OK
8.Connecting to node 127.0.0.1:7003: OK
9.>>> Performing Cluster Check (using node 127.0.0.1:7006)
0.M: efc3131fbdc6cf929720e0ef7136cae85657481 127.0.0.1:7006
1.slots: (0 slots) master
2.0 additional replica(s)
3.M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
4.slots:5461-10922 (5462 slots) master
5.1 additional replica(s)
6.S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
7.slots: (0 slots) slave
8.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
9.S: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
10.slots: (0 slots) slave
11.replicates d2237fdcfbba672de766b913d1186cebc6e1761
12.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
13.slots:10923-16383 (5461 slots) master
14.1 additional replica(s)
15.S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
16.slots: (0 slots) slave
17.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
18.M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
```

```
9.slots:0-5460 (5461 slots) master
```

```
10.1 additional replica(s)
```

```
1.[OK] All nodes agree about slots configuration.
```

```
2.>>> Check for open slots...
```

```
3.>>> Check slots coverage...
```

```
4.[OK] All 16384 slots covered.
```

```
5.How many slots do you want to move (from 1 to 16384)?
```

它提示我们需要迁移多少slot到7006上，我们可以算一下：16384/4 = 4096，也就是说，为了平衡分配起见，我们需要移动4096个槽点到7006上。好，那输入4096:

```
1.How many slots do you want to move (from 1 to 16384)? 4096
```

```
2.What is the receiving node ID?
```

它又提示我们，接受的node ID是多少，7006的id 我们通过上面就可以看到是efc3131fbdc6cf929720e0e0f7136cae85657481:

```
1.What is the receiving node ID? efc3131fbdc6cf929720e0e0f7136cae85657481
```

```
2.Please enter all the source node IDs.
```

```
3.Type &#39;all&#39; to use all the nodes as source nodes for the hash slots.
```

```
4.Type &#39;done&#39; once you entered all the source nodes IDs.
```

```
5.Source node #1:
```

接着，redis-trib 会向你询问重新分片的源节点（source node），也即是，要从哪个节点中取出 4096 个哈希槽，并将这些槽移动到7006节点上面。

如果我们不打算从特定的节点上取出指定数量的哈希槽，那么可以向 redis-trib 输入 all，这样的话，集群中的所有主节点都会成为源节点，redis-trib 将从各个源节点中各取出一部分哈希槽，凑够 4096 个，然后移动到7006节点上:

```
1.Source node #1:all
```

接下来就开始迁移了，并且会询问你是否确认:

```
1.Moving slot 1359 from d2237fdcfbba672de766b913d1186cebc6e1761
```

```
2.Moving slot 1360 from d2237fdcfbba672de766b913d1186cebc6e1761
```

```
3.Moving slot 1361 from d2237fdcfbba672de766b913d1186cebc6e1761
```

```
4.Moving slot 1362 from d2237fdcfbba672de766b913d1186cebc6e1761
```

```
5.Moving slot 1363 from d2237fdcfbba672de766b913d1186cebc6e1761
```

```
6.Moving slot 1364 from d2237fdcfbba672de766b913d1186cebc6e1761
```

```
7.Do you want to proceed with the proposed reshard plan (yes/no)?
```

输入 yes 并使用按下回车之后，redis-trib 就会正式开始执行重新分片操作，将指定的哈希槽从源节点一个个地移动到7006节点上面。

迁移完毕之后，我们来检查下:

```
1.[root@web3 7006]# redis-trib.rb check 127.0.0.1:7000
```

```
2.Connecting to node 127.0.0.1:7000: OK
```

```
3.Connecting to node 127.0.0.1:7006: OK
```

```
4.Connecting to node 127.0.0.1:7004: OK
```

```
5.Connecting to node 127.0.0.1:7005: OK
```

```
6.Connecting to node 127.0.0.1:7003: OK
```

```
7.Connecting to node 127.0.0.1:7001: OK
```

```
8.Connecting to node 127.0.0.1:7002: OK
```

```
9.>>> Performing Cluster Check (using node 127.0.0.1:7000)
```

```
0.S: 3707debcb7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
```

```
1.slots: (0 slots) slave
```

```
2.replicates d2237fdcfbba672de766b913d1186cebc6e1761
```

```
3.M: efc3131fbdc6cf929720e0e0f7136cae85657481 127.0.0.1:7006
```

```
4.slots:0-1364,5461-6826,10923-12287 (4096 slots) master
```

```
5.0 additional replica(s)
```

```
6.S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
```

```
7.slots: (0 slots) slave
```

```
8.replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
```

```
9.S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
```

```
10.slots: (0 slots) slave
```

```
11.replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
```

```
12.M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
```

```
13.slots:1365-5460 (4096 slots) master
```

```
14.1 additional replica(s)
```

```
15.M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
```

```
16.slots:6827-10922 (4096 slots) master
```

```
17.1 additional replica(s)
```

```
18.M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
```

```
19.slots:12288-16383 (4096 slots) master
```

```
20.1 additional replica(s)
```

```
1.[OK] All nodes agree about slots configuration.
```

```
2.>>> Check for open slots...
```

```
3.>>> Check slots coverage...
```

```
4.[OK] All 16384 slots covered.
```

我们着重看7006:

```
0-1364,5461-6826,10923-12287 (4096 slots)
```

这些原来在其他节点上的slot 杯迁移到了7006上。原来，它只是间隔的移动，并不是衔接的整体移动，我们看下有数据了没?

```
1.[root@web3 7006]# redis-cli -c -p 7006
```

```
2.127.0.0.1:7006> keys *
```

```
3.1) "city"
```

```
4. 2) "age"
5. 3) "citylocate"
6. 127.0.0.1:7006> get city
7. "shanghai"
```

非常赞，已经有数据了。

## 2. 新建一个7007从节点，作为7006的从节点

我们再新建一个节点7007，步骤类似，就先省略了。建好后，启动起来，我们看如何把它加入到集群中的从节点中：

```
[root@web3 7007]# redis-trib.rb add-node --slave 127.0.0.1:7007 127.0.0.1:7000
```

add-node的时候加上--slave表示是加入到从节点中，但是这样加，是随机的。这里的命令行完全像我们在添加一个新主服务器时使用的一样，所以我们没有指定要给哪个主服务器添加副本。这种情况下，redis-trib会将7007作为一个具有较少副本的随机的主服务器的副本。

那么，你猜，它会作为谁的从节点，应该是7006，因为7006还没有从节点。我们运行下。

```
1. [root@web3 7007]# redis-trib.rb add-node --slave 127.0.0.1:7007 127.0.0.1:7000
2. ...
3. ...
4. [OK] All 16384 slots covered.
5. Automatically selected master 127.0.0.1:7006
6. Connecting to node 127.0.0.1:7007: OK
7. >>> Send CLUSTER MEET to node 127.0.0.1:7007 to make it join the cluster.
8. Waiting for the cluster to join.
9. >>> Configure node as replica of 127.0.0.1:7006.
10. [OK] New node added correctly.
```

上面提示说，自动选择了7006作为master节点。并且成功了。我们检查下：

```
1. [root@web3 7007]# redis-trib.rb check 127.0.0.1:7000
2. Connecting to node 127.0.0.1:7000: OK
3. Connecting to node 127.0.0.1:7006: OK
4. Connecting to node 127.0.0.1:7004: OK
5. Connecting to node 127.0.0.1:7005: OK
6. Connecting to node 127.0.0.1:7003: OK
7. Connecting to node 127.0.0.1:7001: OK
8. Connecting to node 127.0.0.1:7007: OK
9. Connecting to node 127.0.0.1:7002: OK
10. >>> Performing Cluster Check (using node 127.0.0.1:7000)
1. S: 3707debcbe7be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000
2. slots: (0 slots) slave
3. replicates d2237fdcfbba672de766b913d1186cebc6e1761
4. M: efc3131fbdc6cf929720e0e0f7136cae85657481 127.0.0.1:7006
5. slots:0-1364,5461-6826,10923-12287 (4096 slots) master
6. 1 additional replica(s)
7. S: 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004
8. slots: (0 slots) slave
9. replicates cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
10. S: 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005
11. slots: (0 slots) slave
12. replicates dfa0754c7854a874a6ebd2613b86140ad97701fc
13. M: d2237fdcfbba672de766b913d1186cebc6e1761 127.0.0.1:7003
14. slots:1365-5460 (4096 slots) master
15. 1 additional replica(s)
16. M: cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001
17. slots:6827-10922 (4096 slots) master
18. 1 additional replica(s)
19. S: 86d05e7c2b197dc182b5e71069e791d033cf899e 127.0.0.1:7007
20. slots: (0 slots) slave
21. replicates efc3131fbdc6cf929720e0e0f7136cae85657481
22. M: dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002
23. slots:12288-16383 (4096 slots) master
24. 1 additional replica(s)
25. [OK] All nodes agree about slots configuration.
26. >>> Check for open slots...
27. >>> Check slots coverage...
28. [OK] All 16384 slots covered.
```

果然，7007加入到了7006的从节点当中。

你说，我如果想指定一个主节点行不行？当然可以。我们再建一个7008节点。

```
1. redis-trib.rb add-node --slave --master-id efc3131fbdc6cf929720e0e0f7136cae85657481 127.0.0.1:7008 127.0.0.1:7000
   --master-id表示指定的主节点node id。这里指定的是7006这个主节点。
1. Waiting for the cluster to join.
2. >>> Configure node as replica of 127.0.0.1:7006.
3. [OK] New node added correctly.
4. [root@web3 7008]#
```

提示我们已经作为7006的附属品，也就是加入到7006的从节点来了，照这么说，7006就有2个从节点了，我们看一下：

```
1. redis-cli -c -p 7008 cluster nodes |grep efc3131fbdc6cf929720e0e0f7136cae85657481
```

```
2.
3. 86d05e7c2b197dc182b5e71069e791d033cf899e 127.0.0.1:7007 slave efc3131fbdc6cf929720e0ef7136cae85657481 0 1445089507786 8 connected
4. efc3131fbdc6cf929720e0ef7136cae85657481 127.0.0.1:7006 master - 0 1445089508289 8 connected 0-1364 5461-6826 10923-12287
5. 44321e7d619410dc4e0a8745366610a0d06d2395 127.0.0.1:7008 myself,slave efc3131fbdc6cf929720e0ef7136cae85657481 0 0 0 connected
```

我们过滤了下看结果，果真，7007和7008是7006的从节点了。  
刚好，我们再做一个实验，我把7006的杀掉，看7007和7008谁会变成主节点：

```
1. [root@web3 7008]# ps -ef|grep redis
2. root 11384 1 0 09:56 ? 00:00:16 redis-server *:7001 [cluster]
3. root 11388 1 0 09:56 ? 00:00:16 redis-server *:7002 [cluster]
4. root 11392 1 0 09:56 ? 00:00:16 redis-server *:7003 [cluster]
5. root 11396 1 0 09:56 ? 00:00:15 redis-server *:7004 [cluster]
6. root 11400 1 0 09:56 ? 00:00:15 redis-server *:7005 [cluster]
7. root 12100 1 0 11:01 ? 00:00:11 redis-server *:7000 [cluster]
8. root 12132 1 0 11:28 ? 00:00:11 redis-server *:7006 [cluster]
9. root 12202 1 0 13:14 ? 00:00:02 redis-server *:7007 [cluster]
0. root 12219 1 0 13:39 ? 00:00:00 redis-server *:7008 [cluster]
1. root 12239 8259 0 13:49 pts/0 00:00:00 grep redis
2. [root@web3 7008]# kill 12132
3. [root@web3 7008]# redis-cli -c -p 7008
4. 127.0.0.1:7008> get ss5rtr
5. -> Redirected to slot [1188] located at 127.0.0.1:7007
6. "66"
7. 127.0.0.1:7007> cluster nodes
8. efc3131fbdc6cf929720e0ef7136cae85657481 127.0.0.1:7006 master,fail - 1445089780668 1445089779963 8 disconnected
9. d2237fdcfbba672de766b913d1186cebcb6e1761 127.0.0.1:7003 master - 0 1445089812195 7 connected 1365-5460
0. 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005 slave dfa0754c7854a874a6ebd2613b86140ad97701fc 0 1445089813710 3 connected
1. 86d05e7c2b197dc182b5e71069e791d033cf899e 127.0.0.1:7007 myself,master - 0 0 10 connected 0-1364 5461-6826 10923-12287
2. cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001 master - 0 1445089814214 2 connected 6827-10922
3. 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004 slave cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 0 1445089812701 2 connected
4. 44321e7d619410dc4e0a8745366610a0d06d2395 127.0.0.1:7008 slave 86d05e7c2b197dc182b5e71069e791d033cf899e 0 1445089814214 10 connected
5. 3707debcb67be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000 slave d2237fdcfbba672de766b913d1186cebcb6e1761 0 1445089813204 7 connected
6. dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002 master - 0 1445089813204 3 connected 12288-16383
7. 127.0.0.1:7007>
```

看，7007获得了成为主节点的机会，7008就变成了7007的从节点。  
那么这个时候，重启7006节点，那么他就会变成了一个7007的从节点了。

```
1. 27.0.0.1:7008> cluster nodes
2. 30858dbf483b61b9838d5c1f853a60beaa4e7afd 127.0.0.1:7005 slave dfa0754c7854a874a6ebd2613b86140ad97701fc 0 1445089986148 3 connected
3. 86d05e7c2b197dc182b5e71069e791d033cf899e 127.0.0.1:7007 master - 0 1445089986652 10 connected 0-1364 5461-6826 10923-12287
4. d2237fdcfbba672de766b913d1186cebcb6e1761 127.0.0.1:7003 master - 0 1445089986148 7 connected 1365-5460
5. cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 127.0.0.1:7001 master - 0 1445089987155 2 connected 6827-10922
6. efc3131fbdc6cf929720e0ef7136cae85657481 127.0.0.1:7006 slave 86d05e7c2b197dc182b5e71069e791d033cf899e 0 1445089985644 10 connected
7. 44321e7d619410dc4e0a8745366610a0d06d2395 127.0.0.1:7008 myself,slave 86d05e7c2b197dc182b5e71069e791d033cf899e 0 0 0 connected
8. dfa0754c7854a874a6ebd2613b86140ad97701fc 127.0.0.1:7002 master - 0 1445089986652 3 connected 12288-16383
9. 4b4aef8b48c427a3c903518339d53b6447c58b93 127.0.0.1:7004 slave cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c 0 1445089987660 2 connected
0. 3707debcb67be66d4a1968eaf3a5ffaf4308efa4 127.0.0.1:7000 slave d2237fdcfbba672de766b913d1186cebcb6e1761 0 1445089985644 7 connected
```

```
1. 127.0.0.1:7008>
3. 移除一个主节点
```

有加肯定有减，redis cluster同样支持移除节点功能，同样也是redis-trib.rb的用法：  
redis-trib del-node 127.0.0.1:7000 ``

和新加节点有点不同的是，移除需要节点的node-id。那我们尝试将7007这个主节点移除：

```
1. [root@web3 7006]# redis-trib.rb del-node 127.0.0.1:7000 86d05e7c2b197dc182b5e71069e791d033cf899e
2. >>> Removing node 86d05e7c2b197dc182b5e71069e791d033cf899e from cluster 127.0.0.1:7000
3. Connecting to node 127.0.0.1:7000: OK
4. Connecting to node 127.0.0.1:7006: OK
5. Connecting to node 127.0.0.1:7004: OK
6. Connecting to node 127.0.0.1:7001: OK
7. Connecting to node 127.0.0.1:7003: OK
8. Connecting to node 127.0.0.1:7007: OK
9. Connecting to node 127.0.0.1:7008: OK
0. Connecting to node 127.0.0.1:7005: OK
1. Connecting to node 127.0.0.1:7002: OK
2. [ERR] Node 127.0.0.1:7007 is not empty! Reshard data away and try again.
```

报错了，它提示我们说，由于7007里面已经有数据了，不能被移除，要先将它的数据转移出去。也就是说得重新分片，用上面增加新节点后的分片方式一样，用我们再来一遍：

```
redis-trib.rb reshard 127.0.0.1:7000
```

由于中间太多内容，就省略不重要的，只简单说明下关键的几步：

```
1. M: 86d05e7c2b197dc182b5e71069e791d033cf899e 127.0.0.1:7007
2. slots:0-1364,5461-6826,10923-12287 (4096 slots) master
3.
4. How many slots do you want to move (from 1 to 16384)?
```

提示, 我们要分多少个槽点, 由于7007上有4096个槽点, 所以这里填写4096

1. How many slots do you want to move (from 1 to 16384)? 4096
2. What is the receiving node ID?

提示我们, 需要移动到哪个id上, 那就填7001的吧:

1. What is the receiving node ID? cb5c04b6160c3b7e18cad5d49d8e2987b27e0d6c
2. Please enter all the source node IDs.
3. Type &#39;all&#39; to use all the nodes as source nodes for the hash slots.
4. Type &#39;done&#39; once you entered all the source nodes IDs.
5. Source node #1:

这里就是关键了, 它要我们从哪个节点去转移数据到7001, 因为我们要删除7007的, 所以, 我们就得7007的id了

1. Source node #1:86d05e7c2b197dc182b5e71069e791d033cf899e
2. Source node #2:done
3. Do you want to proceed with the proposed reshard plan (yes/no)? yes

ok,这样就迁移好了。我们看看7007是否为空了:

1. [root@web3 7006]# redis-cli -c -p 7007
2. 127.0.0.1:7007> keys \*
3. (empty list or set)
4. 127.0.0.1:7007> cluster nodes
- 5.
6. 86d05e7c2b197dc182b5e71069e791d033cf899e 127.0.0.1:7007 myself, master - 0 0 10 connected

果然为空了, 好。现在再进行移除节点操作:

1. [root@web3 7006]# redis-trib.rb del-node 127.0.0.1:7000 86d05e7c2b197dc182b5e71069e791d033cf899e
2. >>> Removing node 86d05e7c2b197dc182b5e71069e791d033cf899e from cluster 127.0.0.1:7000
3. Connecting to node 127.0.0.1:7000: OK
4. Connecting to node 127.0.0.1:7006: OK
5. Connecting to node 127.0.0.1:7004: OK
6. Connecting to node 127.0.0.1:7001: OK
7. Connecting to node 127.0.0.1:7003: OK
8. Connecting to node 127.0.0.1:7007: OK
9. Connecting to node 127.0.0.1:7008: OK
0. Connecting to node 127.0.0.1:7005: OK
1. Connecting to node 127.0.0.1:7002: OK
2. >>> Sending CLUSTER FORGET messages to the cluster...
3. >>> 127.0.0.1:7006 as replica of 127.0.0.1:7001
4. >>> 127.0.0.1:7008 as replica of 127.0.0.1:7001
5. >>> SHUTDOWN the node.

哈哈哈哈哈, 删除成功, 而且还很人性化的将7006和70082个没了爹的孩子送给了7001。好评

来源: <http://m.2cto.com/kf/201606/514203.html>